
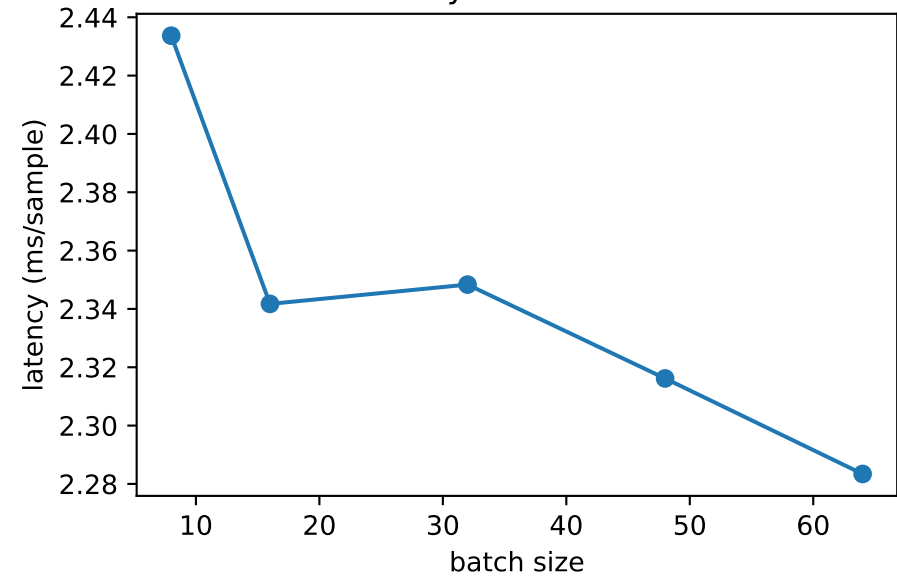


Efficiency & resource curves


Throughput vs. batch size



Latency vs. batch size



Memory vs. batch size



Throughput vs. length (bs=32)

